

# Lab 4

Jia Yu Lin

11:59PM March 11, 2021

Load up the famous iris dataset. We are going to do a different prediction problem. Imagine the only input  $x$  is Species and you are trying to predict  $y$  which is Petal.Length. A reasonable prediction is the average petal length within each Species. Prove that this is the OLS model by fitting an appropriate `lm` and then using the `predict` function to verify.

```
data(iris)
mod =lm( Petal.Length ~ Species, iris)
mod

##
## Call:
## lm(formula = Petal.Length ~ Species, data = iris)
##
## Coefficients:
##      (Intercept)  Speciesversicolor  Speciesvirginica
##              1.462              2.798              4.090

mean(iris$Petal.Length[iris$Species == 'setosa'])

## [1] 1.462

mean(iris$Petal.Length[iris$Species == 'versicolor'])

## [1] 4.26

mean(iris$Petal.Length[iris$Species == 'virginica'])

## [1] 5.552

predict(mod, data.frame(Species = c("setosa")))

##      1
## 1.462

predict(mod, data.frame(Species = c("versicolor")))

##      1
## 4.26
```

```
predict(mod, data.frame(Species = c("virginica")))
```

```
##      1  
## 5.552
```

Construct the design matrix with an intercept,  $X$ , without using `model.matrix`.

```
#TO-DO  
X = cbind(1, iris$Species == 'versicolor', iris$Species == 'virginica')  
View(X)  
head(X)
```

```
##      [,1] [,2] [,3]  
## [1,]    1    0    0  
## [2,]    1    0    0  
## [3,]    1    0    0  
## [4,]    1    0    0  
## [5,]    1    0    0  
## [6,]    1    0    0
```

Find the hat matrix  $H$  for this regression.

```
#TO-DO  
H = X %*% solve(t(X) %*% X) %*% t(X)  
Matrix::rankMatrix(H)
```

```
## [1] 3  
## attr("method")  
## [1] "tolNorm2"  
## attr("useGrad")  
## [1] FALSE  
## attr("tol")  
## [1] 3.330669e-14
```

```
#head(H)
```

Verify this hat matrix is symmetric using the `expect_equal` function in the package `testthat`.

```
#TO-DO  
pacman::p_load(testthat)  
expect_equal(H, t(H))
```

Verify this hat matrix is idempotent using the `expect_equal` function in the package `testthat`.

```
#TO-DO  
expect_equal(H, H%*%H)
```

Using the `diag` function, find the trace of the hat matrix.

```
#TO-DO
#trace is the sum of diagonal
sum(diag(H))
```

```
## [1] 3
```

```
#sum of trace is rank
```

It turns out the trace of a hat matrix is the same as its rank! But we don't have time to prove these interesting and useful facts..

For masters students: create a matrix  $X_{\perp}$ .

```
#TO-DO
```

Using the hat matrix, compute the  $\hat{y}$  vector and using the projection onto the residual space, compute the  $e$  vector and verify they are orthogonal to each other.

```
#TO-DO
y = iris$Petal.Length
y_hat = H %*% iris$Petal.Length
#we are supposed to see y bars for setosa, versicolor, or virginica
#table(y_hat)
I = diag(nrow(iris))
e = (I - H) %*% y
head(e)
```

```
##      [,1]
## [1,] -0.062
## [2,] -0.062
## [3,] -0.162
## [4,]  0.038
## [5,] -0.062
## [6,]  0.238
```

```
Matrix::rankMatrix(I-H)
```

```
## [1] 147
## attr(,"method")
## [1] "tolNorm2"
## attr(,"useGrad")
## [1] FALSE
## attr(,"tol")
## [1] 3.330669e-14
```

Compute SST, SSR and SSE and  $R^2$  and then show that  $SST = SSR + SSE$ .

```
#TO-DO
SSE = t(e) %*% e #same thing is sum(e^2)
y_bar = mean(y)
SST = t(y - y_bar) %*% (y - y_bar)

Rsqr = 1 - SSE/SST
Rsqr
```

```
##           [,1]
## [1,] 0.9413717
```

```
SSR = t(y_hat - y_bar) %*% (y_hat - y_bar)
SSR
```

```
##           [,1]
## [1,] 437.1028
```

```
expect_equal(SSE+SSR, SST)
#this would mean each species would have similar petal length
```

Find the angle  $\theta$  between  $y - \bar{y}1$  and  $\hat{y} - \bar{y}1$  and then verify that its cosine squared is the same as the  $R^2$  from the previous problem.

```
#TO-DO
#231 formula
theta = acos(t(y - y_bar) %*% (y_hat - y_bar) / sqrt(SST * SSR))
#Rsq was pretty large so theta should be pretty small
#theta is 14 degrees
theta * (180/pi)
```

```
##           [,1]
## [1,] 14.01245
```

```
cos(theta)^2
```

```
##           [,1]
## [1,] 0.9413717
```

```
expect_equal(cos(theta)^2, Rsq)
```

Project the  $y$  vector onto each column of the  $X$  matrix and test if the sum of these projections is the same as  $\hat{y}$ .

```
#TO-DO
proj1= (X[,1] %*% t(X[,1]) / as.numeric(t(X[,1]) %*% X[,1])) %*% y #H on to X1
proj2= (X[,2] %*% t(X[,2]) / as.numeric(t(X[,2]) %*% X[,2])) %*% y #H on to X2
proj3= (X[,3] %*% t(X[,3]) / as.numeric(t(X[,3]) %*% X[,3])) %*% y #H on to X3

expect_equal(proj1+proj2+proj3, y_hat)
#this will fail which is what we want
```

Construct the design matrix without an intercept,  $X$ , without using `model.matrix`.

```
#TO-DO
X = cbind( 1e-4, as.numeric(iris$Species == 'versicolor'), as.numeric(iris$Species == 'virginica'))
head(X)
```

```
##      [,1] [,2] [,3]
## [1,] 1e-04  0    0
## [2,] 1e-04  0    0
## [3,] 1e-04  0    0
## [4,] 1e-04  0    0
## [5,] 1e-04  0    0
## [6,] 1e-04  0    0
```

```
#iris
```

Find the OLS estimates using this design matrix. It should be the sample averages of the petal lengths within species.

```
#TO-DO
```

```
mod_X = lm(Petal.Length ~ X, iris)
mod_X
```

```
##
## Call:
## lm(formula = Petal.Length ~ X, data = iris)
##
## Coefficients:
## (Intercept)      X1      X2      X3
##      1.462      NA    2.798    4.090
```

Verify the hat matrix constructed from this design matrix is the same as the hat matrix constructed from the design matrix with the intercept. (Fact: orthogonal projection matrices are unique).

```
#TO-DO
```

```
H_new = X %*% solve(t(X) %*% X) %*% t(X)
expect_equal(H_new, H, tol= 1e-4)
```

Project the  $y$  vector onto each column of the  $X$  matrix and test if the sum of these projections is the same as  $\hat{y}$ .

```
#TO-DO
```

```
proj1= (X[,1] %*% t(X[,1]) / as.numeric(t(X[,1]) %*% X[,1])) %*% y #H on to X1
proj2= (X[,2] %*% t(X[,2]) / as.numeric(t(X[,2]) %*% X[,2])) %*% y #H on to X2
proj3= (X[,3] %*% t(X[,3]) / as.numeric(t(X[,3]) %*% X[,3])) %*% y #H on to X3
expect_equal(proj1+proj2+proj3, y_hat)
```

Convert this design matrix into  $Q$ , an orthonormal matrix.

```
#TO-DO
```

```
qrX = qr(X)
Q = qr.Q(qrX)
```

Project the  $y$  vector onto each column of the  $Q$  matrix and test if the sum of these projections is the same as  $\hat{y}$ .

*#TO-DO*

```
proj1= (Q[,1] %*% t(Q[,1]) / as.numeric(t(Q[,1]) %*% Q[,1])) %*% y #H on to X1
proj2= (Q[,2] %*% t(Q[,2]) / as.numeric(t(Q[,2]) %*% Q[,2])) %*% y #H on to X2
proj3= (Q[,3] %*% t(Q[,3]) / as.numeric(t(Q[,3]) %*% Q[,3])) %*% y #H on to X3

expect_equal(proj1+proj2+proj3, y_hat)
```

Find the  $p = 3$  linear OLS estimates if  $Q$  is used as the design matrix using the `lm` method. Is the OLS solution the same as the OLS solution for  $X$ ?

*#TO-DO*

```
mod_Q = lm(Petal.Length ~ Q, iris)
mod_Q

##
## Call:
## lm(formula = Petal.Length ~ Q, data = iris)
##
## Coefficients:
## (Intercept)      Q1      Q2      Q3
##      3.758      NA    4.347   20.450
```

*#It is not the same as the OLS solution for X*

Use the `predict` function and ensure that the predicted values are the same for both linear models: the one created with  $X$  as its design matrix and the one created with  $Q$  as its design matrix.

*#TO-DO*

```
pred_X = predict(mod_X)
pred_Q = predict(mod_Q)

expect_equal(pred_X, pred_Q)
```

Clear the workspace and load the boston housing data and extract  $X$  and  $y$ . The dimensions are  $n = 506$  and  $p = 13$ . Create a matrix that is  $(p + 1) \times (p + 1)$  full of NA's. Label the columns the same columns as  $X$ . Do not label the rows. For the first row, find the OLS estimate of the  $y$  regressed on the first column only and put that in the first entry. For the second row, find the OLS estimates of the  $y$  regressed on the first and second columns of  $X$  only and put them in the first and second entries. For the third row, find the OLS estimates of the  $y$  regressed on the first, second and third columns of  $X$  only and put them in the first, second and third entries, etc. For the last row, fill it with the full OLS estimates.

*#TO-DO*

```
rm(list = ls())

boston = MASS::Boston
##?Boston
X = cbind(1,as.matrix(boston[, 1:13]))
Y = boston[, 14]
```

```
X_matrix = matrix(NA, nrow= 14, ncol = 14)

colnames(X_matrix) <- c(colnames(X))
X_matrix
```

```
##          crim zn indus chas nox rm age dis rad tax ptratio black lstat
## [1,] NA    NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [2,] NA    NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [3,] NA    NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [4,] NA    NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [5,] NA    NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [6,] NA    NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [7,] NA    NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [8,] NA    NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [9,] NA    NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [10,] NA   NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [11,] NA   NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [12,] NA   NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [13,] NA   NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
## [14,] NA   NA NA    NA    NA NA NA NA NA NA NA    NA    NA    NA
```

```
for( i in 1:ncol(X)){
  mod_coef = coef(lm(Y ~ X[,1:i], data= as.data.frame(boston)))
  count = 1
  for( j in 2:(length(mod_coef))){
    X_matrix[i, count] = mod_coef[j]
    X_matrix[i, 1] = mod_coef[1]
    count = count+1
  }
}
X_matrix
```

```
##          crim          zn          indus          chas          nox
## [1,] 22.5328063          NA          NA          NA          NA
## [2,] 24.0331062 -0.4151903          NA          NA          NA
## [3,] 22.4856281 -0.3520783 0.11610909          NA          NA
## [4,] 27.3946468 -0.2486283 0.05850082 -0.41557782          NA
## [5,] 27.1128031 -0.2287981 0.05928665 -0.44032511 6.894059          NA
## [6,] 29.4899406 -0.2185190 0.05511047 -0.38348055 7.026223 -5.424659
## [7,] -17.9546350 -0.1769135 0.02128135 -0.14365267 4.784684 -7.184892
## [8,] -18.2649261 -0.1727607 0.01421402 -0.13089918 4.840730 -4.357411
## [9,] 0.8274820 -0.1977868 0.06099257 -0.22573089 4.577598 -14.451531
## [10,] 0.1553915 -0.1780398 0.06095248 -0.21004328 4.536648 -13.342666
## [11,] 2.9907868 -0.1795543 0.07145574 -0.10437742 4.110667 -12.591596
## [12,] 27.1523679 -0.1840321 0.03909990 -0.04232450 3.487528 -22.182110
## [13,] 20.6526280 -0.1599391 0.03887365 -0.02792186 3.216569 -20.484560
## [14,] 36.4594884 -0.1080114 0.04642046 0.02055863 2.686734 -17.766611
##          rm          age          dis          rad          tax          ptratio
## [1,]          NA          NA          NA          NA          NA          NA
## [2,]          NA          NA          NA          NA          NA          NA
## [3,]          NA          NA          NA          NA          NA          NA
## [4,]          NA          NA          NA          NA          NA          NA
```

```
## [5,]      NA      NA      NA      NA      NA      NA
## [6,]      NA      NA      NA      NA      NA      NA
## [7,]  7.341586      NA      NA      NA      NA      NA
## [8,]  7.386357 -0.0236248493      NA      NA      NA      NA
## [9,]  6.752352 -0.0556354540 -1.760312      NA      NA      NA
## [10,] 6.791184 -0.0562612189 -1.748296 -0.04529059      NA      NA
## [11,] 6.664084 -0.0546675064 -1.727933  0.15926305 -0.01434060      NA
## [12,] 6.075744 -0.0451880522 -1.583852  0.25472196 -0.01221262 -0.9962062
## [13,] 6.123072 -0.0459320518 -1.554912  0.28157503 -0.01173838 -1.0142228
## [14,] 3.809865  0.0006922246 -1.475567  0.30604948 -0.01233459 -0.9527472
##          black      lstat
## [1,]      NA      NA
## [2,]      NA      NA
## [3,]      NA      NA
## [4,]      NA      NA
## [5,]      NA      NA
## [6,]      NA      NA
## [7,]      NA      NA
## [8,]      NA      NA
## [9,]      NA      NA
## [10,]      NA      NA
## [11,]      NA      NA
## [12,]      NA      NA
## [13,] 0.013620833      NA
## [14,] 0.009311683 -0.5247584
```

Why are the estimates changing from row to row as you add in more predictors?

#TO-DO The estimates are changing as more predictors are added because it is trying to find a better fit.

Create a vector of length  $p + 1$  and compute the  $R^2$  values for each of the above models.

```
#TO-DO
y <- c()
for( j in 1:14){
  mod_coef = lm(Y~X[,1:j])
  y <- append(y, summary(mod_coef)$r.squared)
}
y
```

```
## [1] 0.0000000 0.1507805 0.2339884 0.2937136 0.3295277 0.3313127 0.5873770
## [8] 0.5894902 0.6311488 0.6319479 0.6396628 0.6703141 0.6842043 0.7406427
```

Is  $R^2$  monotonically increasing? Why?

#TO-DO R square increases because the amount of features we have went up